

Tree-Based Methods

Classification Trees

November 3, 2017

Classification Trees

- A classification tree is very similar to a regression tree, except that it is used to predict a qualitative response rather than a quantitative one.
- Recall that for a regression tree, the predicted response for an observation is given by the mean response of the training observations that belong to the same terminal node.
- In contrast, for a classification tree, we predict that each observation belongs to the most commonly occurring class of training observations in the region to which it belongs.
- In interpreting the results of a classification tree, we are often interested not only in the class prediction corresponding to a particular terminal node region, but also in the class proportions among the training observations that fall into that region.
- However, in the classification setting, RSS cannot be used as a criterion for making the binary splits.
- A natural alternative to RSS is the *classification error rate*.
- The classification error rate is simply the fraction of the training observations in that region that do not belong to the most common class.




Figure 1: The unpruned tree.

Figure 1 shows an example on the **Heart data** set. These data contain a binary outcome HD for 303 patients who presented with chest pain. An outcome value of Yes indicates the presence of heart disease based on an angiographic test, while No means no heart disease. There are 13 predictors including Age, Sex, Chol (a cholesterol measurement), and other heart and lung function measurements.




Figure 2: Left: Cross -validation error, training, and test error, for different sizes of the pruned tree. Bottom Right: The pruned tree corresponding to the minimal cross-validation error.

Fitting Classification Trees

The **tree library** is used to construct classification and regression trees.

Creat the data set

We first use classification trees to analyze the Carseats data set. In these data, Sales is a continuous variable, and so we begin by recoding it as a binary variable. We use the **ifelse()** function to create a variable, called High, which takes on a value of Yes if the Sales variable exceeds 8, and takes on a value of No otherwise.

```
> library (ISLR)
> attach (Carseats )
> High=ifelse (Sales <=8," No"," Yes ")
> Carseats =data.frame(Carseats ,High)
```

Finally, we use the `data.frame()` function to merge High with the rest of the Carseats data.

We now use the **tree()** function to fit a classification tree in order to predict High using all variables but Sales.

```
> tree.carseats =tree(High.-Sales ,Carseats ) <---- Using All Data
> summary (tree.carseats )
```

Classification tree:

```
tree(formula = High ~ . - Sales, data = Carseats)
```

Variables actually used in tree construction:

```
[1] "ShelveLoc" "Price" "US" "Income" "CompPrice" "Population" "Adver"
```

[8] "Age"

Number of terminal nodes: 27

Residual mean deviance: 0.4575 = 170.7 / 373

Misclassification error rate: 0.09 = 36 / 400

The `summary()` function lists the variables that are used as internal nodes in the tree, the number of terminal nodes, and the (training) error rate. We see that the training error rate is 9 %.

A **small deviance** indicates a tree that provides a good fit to the (training) data.

One of the most attractive properties of trees is that they can be graphically displayed. We use the `plot()` function to display the tree structure, and the `text()` function to display the node labels. The argument `pretty=0` instructs R to include the category names for any qualitative predictors, rather than simply displaying a letter for each category.

```
> plot(tree.carseats )
> text(tree.carseats ,pretty =0)
```




Figure 3: Decision Tree

If we just type the name of the tree object, R prints output corresponding to each branch of the tree. R displays the split criterion (e.g. `Price < 92.5`), the number of observations in that branch, the deviance, the overall prediction for the branch (Yes or No), and the fraction of observations in that branch that take on values of Yes and No. Branches that lead to terminal nodes are indicated using asterisks.

```
> tree.carseats
node), split, n, deviance, yval, (yprob)
* denotes terminal node
```

```
1) root 400 541.500 No ( 0.41000 0.59000 )
2) ShelveLoc: Good 85 90.330 Yes ( 0.77647 0.22353 )
4) Price < 135 68 49.260 Yes ( 0.88235 0.11765 )
```

```
8) US: No 17 22.070 Yes ( 0.64706 0.35294 )
16) Price < 109 8 0.000 Yes ( 1.00000 0.00000 ) *
```

Properly evaluate the performance of a classification tree

In order to properly evaluate the performance of a classification tree on these data, we must estimate the test error rather than simply computing the training error. We split the observations into a training set and a test set, build the tree using the training set, and evaluate its performance on the test data. The `predict()` function can be used for this purpose. In the case of a classification tree, the argument `type="class"` instructs R to return the actual class prediction.

```
> set.seed (2)
> train=sample (1: nrow(Carseats ), 200)
> Carseats .test=Carseats [-train ,]
> High.test=High[-train ]
> tree.carseats =tree(High.-Sales ,Carseats ,subset =train )
> tree.pred=predict (tree.carseats ,Carseats .test ,type =" class ")
> table(tree.pred ,High.test)
      High.test
tree.pred No  Yes
      No  86  27
      Yes 30  57
> (86+57) /200
[1] 0.715
```

This approach leads to correct predictions for around 71.5% of the locations in the test data set. Next, we consider whether pruning the tree might lead to improved results. The function `cv.tree()` performs cross-validation in order to `cv.tree()` determine the optimal level of tree complexity; cost complexity pruning is used in order to select a sequence of trees for consideration. We use the argument `FUN=prune.misclass` in order to indicate that we want the classification error rate to guide the cross-validation and pruning process, rather than the default for the `cv.tree()` function, which is deviance. The `cv.tree()` function reports the number of terminal nodes of each tree considered (size) as well as the corresponding error rate and the value of the cost-complexity parameter used.

```
> set.seed (3)
> cv.carseats =cv.tree(tree.carseats ,FUN=prune.misclass )
> names(cv.carseats )
[1] "size" "dev " "k" "method "
> cv.carseats
$size
[1] 19 17 14 13 9 7 3 2 1
$dev
[1] 55 55 53 52 50 56 69 65 80
$k
[1] -Inf 0.0000000 0.6666667 1.0000000 1.7500000
2.0000000 4.2500000
[8] 5.0000000 23.0000000
$method
[1] "misclass "
```

```
attr(," class ")
[1] "prune" "tree.sequence "
```

Note that, despite the name, **dev** corresponds to the **cross-validation error** rate in this instance. The tree with 9 terminal nodes results in the lowest cross-validation error rate, with 50 cross-validation errors. We plot the error rate as a function of both *size* and *k*.

```
> par(mfrow =c(1,2))
> plot(cv.carseats$size ,cv.carseats$dev ,type="b")
> plot(cv.carseats$k ,cv.carseats$dev ,type="b")
```




Figure 4

We now apply the **prune.misclass()** function in order to prune the tree to obtain the nine-node tree.

```
> prune.carseats =prune.misclass (tree.carseats ,best =9)
> plot(prune.carseats )
> text(prune.carseats ,pretty =0)
```

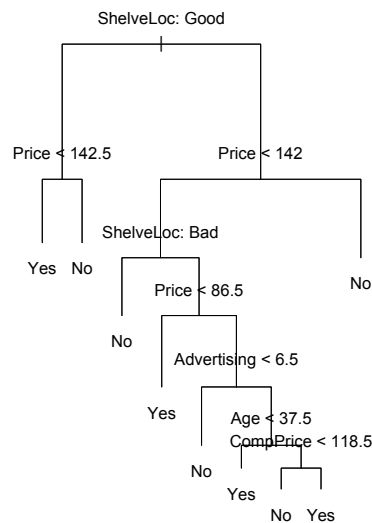


Figure 5

How well does this pruned tree perform on the test data set? Once again, we apply the **predict()** function.

```

> tree.pred=predict (prune.carseats , Carseats .test ,type=" class ")
> table(tree.pred ,High.test)
      High.test
tree.pred No   Yes
      No   94   24
      Yes   22   60
> (94+60) /200
[1] 0.77

```

Now 77% of the test observations are correctly classified, so not only has the pruning process produced a more interpretable tree, but it has also improved the classification accuracy.

R Code

```

library (tree)
library (MASS)
library (ISLR)
?Carseats
View(Carseats)

```

```

##### Creat the proper data
attach (Carseats )
High=ifelse (Sales <=8,"No"," Yes ")
Carseats =data.frame(Carseats ,High)
##### Fit the Tree

```

```
tree.carseats =tree(High~.-Sales ,Carseats )
summary (tree.carseats )
##### plot
plot(tree.carseats )
text(tree.carseats ,pretty =0)
#####
tree.carseats
##### Using Training and Test Data
set.seed(2)
train=sample (1: nrow(Carseats ), 200)
Carseats.test=Carseats [-train ,]
High.test=High[-train ]
#####
tree.carseats =tree(High~.-Sales ,Carseats ,subset =train ) ###<--- Using Training data to bu
tree.pred=predict (tree.carseats ,Carseats.test ,type ="class")##<--- Using Testing data to pre
table(tree.pred ,High.test)
(86+57) /200
#####Prunning the Tree
set.seed(3)
cv.carseats =cv.tree(tree.carseats ,FUN=prune.misclass )
names(cv.carseats )
cv.carseats
#####
par(mfrow =c(1,2))
plot(cv.carseats$size ,cv.carseats$dev ,type="b")
plot(cv.carseats$k ,cv.carseats$dev ,type="b")
#####Plot the Pruned Tree withn ni
prune.carseats =prune.misclass (tree.carseats ,best=9)
plot(prune.carseats )
text(prune.carseats ,pretty =0)
#####How well this mosel predict?
tree.pred=predict (prune.carseats , Carseats.test ,type="class")
table(tree.pred ,High.test)
(94+60) /200
```